

# Eos Compilers



**Fernanda Foertter**  
**HPC User Assistance Specialist**



U.S. DEPARTMENT OF  
**ENERGY**



**OAK RIDGE NATIONAL LABORATORY**

MANAGED BY UT-BATTELLE FOR THE DEPARTMENT OF ENERGY

# Available Compilers

- Cray, Intel, GNU and PGI
- Cray compiler wrappers **Always use the wrapper!**
  - Same command (ftn/cc/CC) regardless of compiler
  - Links in MPI/Libsci libraries automatically
  - Intel is the default
  - Static linking is default

```
module swap PrgEnv-cray PrgEnv-intel
```

# Compiler tips

- Some are more aggressive than others at optimization and standards
- Safe to assume vendor knows its architecture best (e.g. Intel) and thus better at optimizing for their own
- No compiler is superior for all codes
- Try different compilers, good way of catching bugs
- Test your application against lower optimization levels and included tests
- Use system provided libraries

# Caution

## Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

# Intel specific tips

| Pros  | Cons  |
|---|---|
| <ul style="list-style-type: none"><li>• Compatibility with Sandybrige</li><li>• Scalar optimization</li><li>• Vectorization</li><li>• Intel MKL</li><li>• Intel specific OpenMP extensions</li><li>• AVX (speeds up vector + FP)</li><li>• OpenMP 4.0 support</li></ul> | <ul style="list-style-type: none"><li>• <del>No Cray math libraries available</del></li></ul> |

# Some Intel Compiler Options

| Flag          | Description  |
|---------------|--|
| none          | Optimizes for speed; no flags, approx -O2 level  |
| -fast         | Maximizes speed across the entire program. Aggressive. Includes -ipo -O3 -no-prec-div -static -xHost |
| -fast -no-ipo | No IPO   |
| -O0           | Disables optimizations   |
| -O1           | Optimizes for speed but keeps binary small   |
| -O2           | Speed + vectorization, inlining  |
| -O3           | O2 + aggressive loop transformations   |
| -mkl=cluster  | Single threaded MKL (BLAS/LAPACK/scaLAPACK/FFTW)   |
| -mkl=parallel | Multi-Threaded MKL (No ScaLAPACK)  |
| -xAVX         | Default on Eos   |

# MKL link line builder

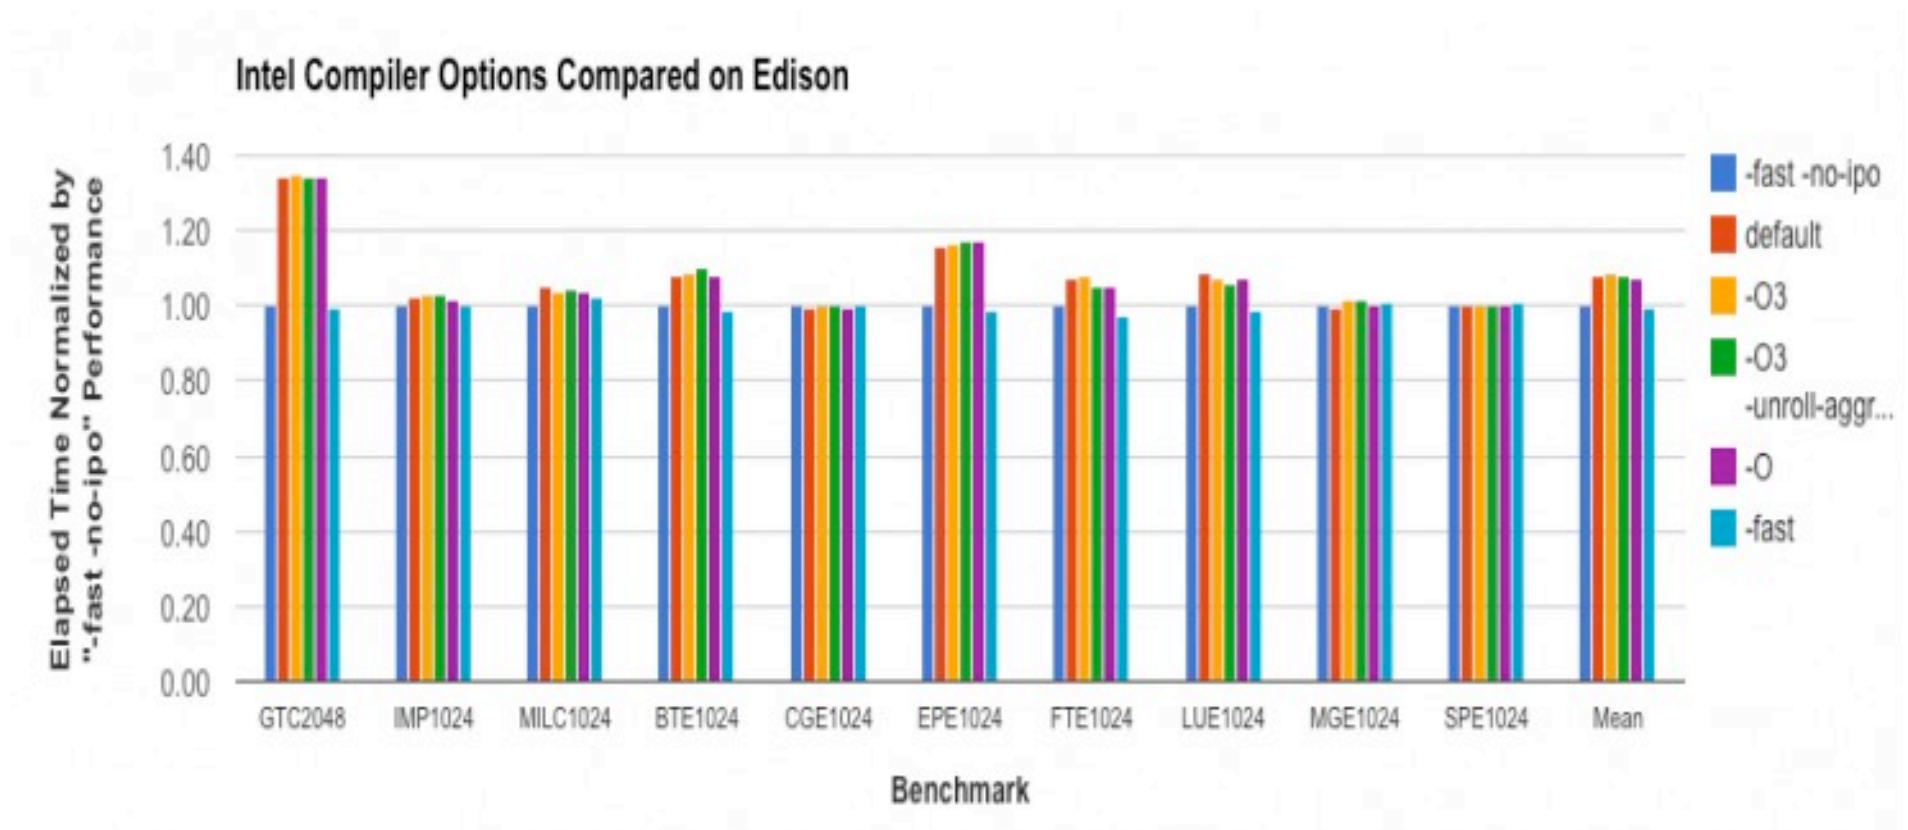
## Intel® Math Kernel Library (Intel® MKL) Link Line Advisor v4.0

|  |   |
|--|---|
| <input type="button" value="Reset"/>                                 |   |
| Select Intel® product:   | Intel(R) MKL 11.1   |
| Select OS:   | <Select operating system>   |
| Select usage model of Intel® Xeon Phi™ Coprocessor:                  | <Select usage model>  |
| Select compiler:   | <Select compiler>   |
| Select architecture:   | <Select architecture>   |
| Select dynamic or static linking:                                    | <Select linking>  |
| Select interface layer:  | <Select interface>  |
| Select sequential or multi-threaded layer:                           | <Select threading>  |
| Select OpenMP library:   | <Select OpenMP>   |
| Select cluster library:  | <input type="checkbox"/> CDFT (BLACS required)<br><input type="checkbox"/> ScaLAPACK (BLACS required)<br><input type="checkbox"/> BLACS |
| Select MPI library:  | <Select MPI>  |
| Select the Fortran 95 interfaces:                                    | <input type="checkbox"/> BLAS95<br><input type="checkbox"/> LAPACK95  |
| Link with Intel® MKL libraries explicitly:                           | <input type="checkbox"/>  |
| Use this link line:<br><Please select all required parameters above> |   |

Source: <http://software.intel.com/en-us/articles/intel-mkl-link-line-advisor>

# Edison Benchmark Results at NERSC

Not a lot of difference (lower is better)




Source: <http://www.nersc.gov/users/computational-systems/edison/performance-and-optimization/compiler-comparisons/>



# Edison Benchmark Results at NERSC

Not a lot of difference (lower is better)



Source: <http://www.nersc.gov/users/computational-systems/edison/performance-and-optimization/compiler-comparisons/>

# Known Issue: Intel OpenMP Hybrid

- The Cray thread affinity settings and Intel's run time OpenMP environment conflict.
- Intel has an extra thread at run time, so 2 threads are scheduled on the same core and the job takes twice as long as it should.

Source: Introducing the Intel Compiler on Edison, Michael Stuart, NUG Feb 2013

# Intel OpenMP Hybrid

- CPU-binding should be turned off
  - Allows user compute threads to spread out over available resources
  - Helper thread will no longer impact performance

## NERSC's Edison Solution

- `OMP_NUM_THREADS <=8`
  - `set KMP_AFFINITY compact + aprun -cc numa_node`
- `OMP_NUM_THREADS >8 and <=16`
  - `set KMP_AFFINITY scatter + aprun -cc none`

Source: Introducing the Intel Compiler on Edison, Michael Stuart, NUG Feb 2013

**OLCF.ORNL.GOV**